

```
In [1]: pip install numpy
```

Collecting numpy
Note: you may need to restart the kernel to use updated packages.

Downloading numpy-1.24.3-cp311-cp311-win_amd64.whl (14.8 MB)
----- 14.8/14.8 MB 789.5 kB/s eta 0:0
0:00
Installing collected packages: numpy
Successfully installed numpy-1.24.3

[notice] A new release of pip available: 22.3.1 -> 23.1.2
[notice] To update, run: python.exe -m pip install --upgrade pip

```
In [2]: pip install pandas
```

Collecting pandas
Downloading pandas-2.0.1-cp311-cp311-win_amd64.whl (10.6 MB)
----- 10.6/10.6 MB 849.5 kB/s eta 0:0
0:00
Requirement already satisfied: python-dateutil>=2.8.2 in c:\users\venka\appdata\local\programs\python\python311\lib\site-packages (from pandas) (2.8.2)
Collecting pytz>=2020.1
Downloading pytz-2023.3-py2.py3-none-any.whl (502 kB)
----- 502.3/502.3 kB 955.5 kB/s eta 0:0
0:00
Collecting tzdata>=2022.1
Downloading tzdata-2023.3-py2.py3-none-any.whl (341 kB)
----- 341.8/341.8 kB 559.2 kB/s eta 0:0
0:00
Requirement already satisfied: numpy>=1.21.0 in c:\users\venka\appdata\local\programs\python\python311\lib\site-packages (from pandas) (1.24.3)
Requirement already satisfied: six>=1.5 in c:\users\venka\appdata\local\programs\python\python311\lib\site-packages (from python-dateutil>=2.8.2->pandas) (1.16.0)
Installing collected packages: pytz, tzdata, pandas
Successfully installed pandas-2.0.1 pytz-2023.3 tzdata-2023.3
Note: you may need to restart the kernel to use updated packages.

[notice] A new release of pip available: 22.3.1 -> 23.1.2
[notice] To update, run: python.exe -m pip install --upgrade pip

```
In [58]: import numpy as np  
import pandas as pd  
import seaborn as sns  
import matplotlib.pyplot as plt
```

```
In [59]: df=pd.read_csv(r"C:\Users\venka\OneDrive\Documents\fiat500_VehicleSelection_D
print(df)
```

	ID	model	engine_power	age_in_days	km	previous_owners
0	1	lounge	51	882	25000	1 \
1	2	pop	51	1186	32500	1
2	3	sport	74	4658	142228	1
3	4	lounge	51	2739	160000	1
4	5	pop	73	3074	106880	1
...
1533	1534	sport	51	3712	115280	1
1534	1535	lounge	74	3835	112000	1
1535	1536	pop	51	2223	60457	1
1536	1537	lounge	51	2557	80750	1
1537	1538	pop	51	1766	54276	1

	lat	lon	price
0	44.907242	8.611560	8900
1	45.666359	12.241890	8800
2	45.503300	11.417840	4200
3	40.633171	17.634609	6000
4	41.903221	12.495650	5700
...
1533	45.069679	7.704920	5200
1534	45.845692	8.666870	4600
1535	45.481541	9.413480	7500
1536	45.000702	7.682270	5990
1537	40.323410	17.568270	7900

[1538 rows x 9 columns]

```
In [60]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1538 entries, 0 to 1537
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  -
0   ID              1538 non-null   int64
1   model           1538 non-null   object
2   engine_power    1538 non-null   int64
3   age_in_days     1538 non-null   int64
4   km              1538 non-null   int64
5   previous_owners 1538 non-null   int64
6   lat             1538 non-null   float64
7   lon             1538 non-null   float64
8   price           1538 non-null   int64
dtypes: float64(2), int64(6), object(1)
memory usage: 108.3+ KB
```

In [61]: `df.head(10)`

Out[61]:

	ID	model	engine_power	age_in_days	km	previous_owners	lat	lon	pri
0	1	lounge	51	882	25000	1	44.907242	8.611560	89
1	2	pop	51	1186	32500	1	45.666359	12.241890	88
2	3	sport	74	4658	142228	1	45.503300	11.417840	42
3	4	lounge	51	2739	160000	1	40.633171	17.634609	60
4	5	pop	73	3074	106880	1	41.903221	12.495650	57
5	6	pop	74	3623	70225	1	45.000702	7.682270	79
6	7	lounge	51	731	11600	1	44.907242	8.611560	107
7	8	lounge	51	1521	49076	1	41.903221	12.495650	91
8	9	sport	73	4049	76000	1	45.548000	11.549470	56
9	10	sport	51	3653	89000	1	45.438301	10.991700	60

In [62]: `df.tail()`

Out[62]:

	ID	model	engine_power	age_in_days	km	previous_owners	lat	lon
1533	1534	sport	51	3712	115280	1	45.069679	7.70492
1534	1535	lounge	74	3835	112000	1	45.845692	8.66687
1535	1536	pop	51	2223	60457	1	45.481541	9.41348
1536	1537	lounge	51	2557	80750	1	45.000702	7.68227
1537	1538	pop	51	1766	54276	1	40.323410	17.56827

In [63]: `df.describe()`

Out[63]:

	ID	engine_power	age_in_days	km	previous_owners	lat
count	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000
mean	769.500000	51.904421	1650.980494	53396.011704	1.123537	43.541361
std	444.126671	3.988023	1289.522278	40046.830723	0.416423	2.133518
min	1.000000	51.000000	366.000000	1232.000000	1.000000	36.855839
25%	385.250000	51.000000	670.000000	20006.250000	1.000000	41.802990
50%	769.500000	51.000000	1035.000000	39031.000000	1.000000	44.394096
75%	1153.750000	51.000000	2616.000000	79667.750000	1.000000	45.467960
max	1538.000000	77.000000	4658.000000	235000.000000	4.000000	46.795612

In [64]: `df.info()`

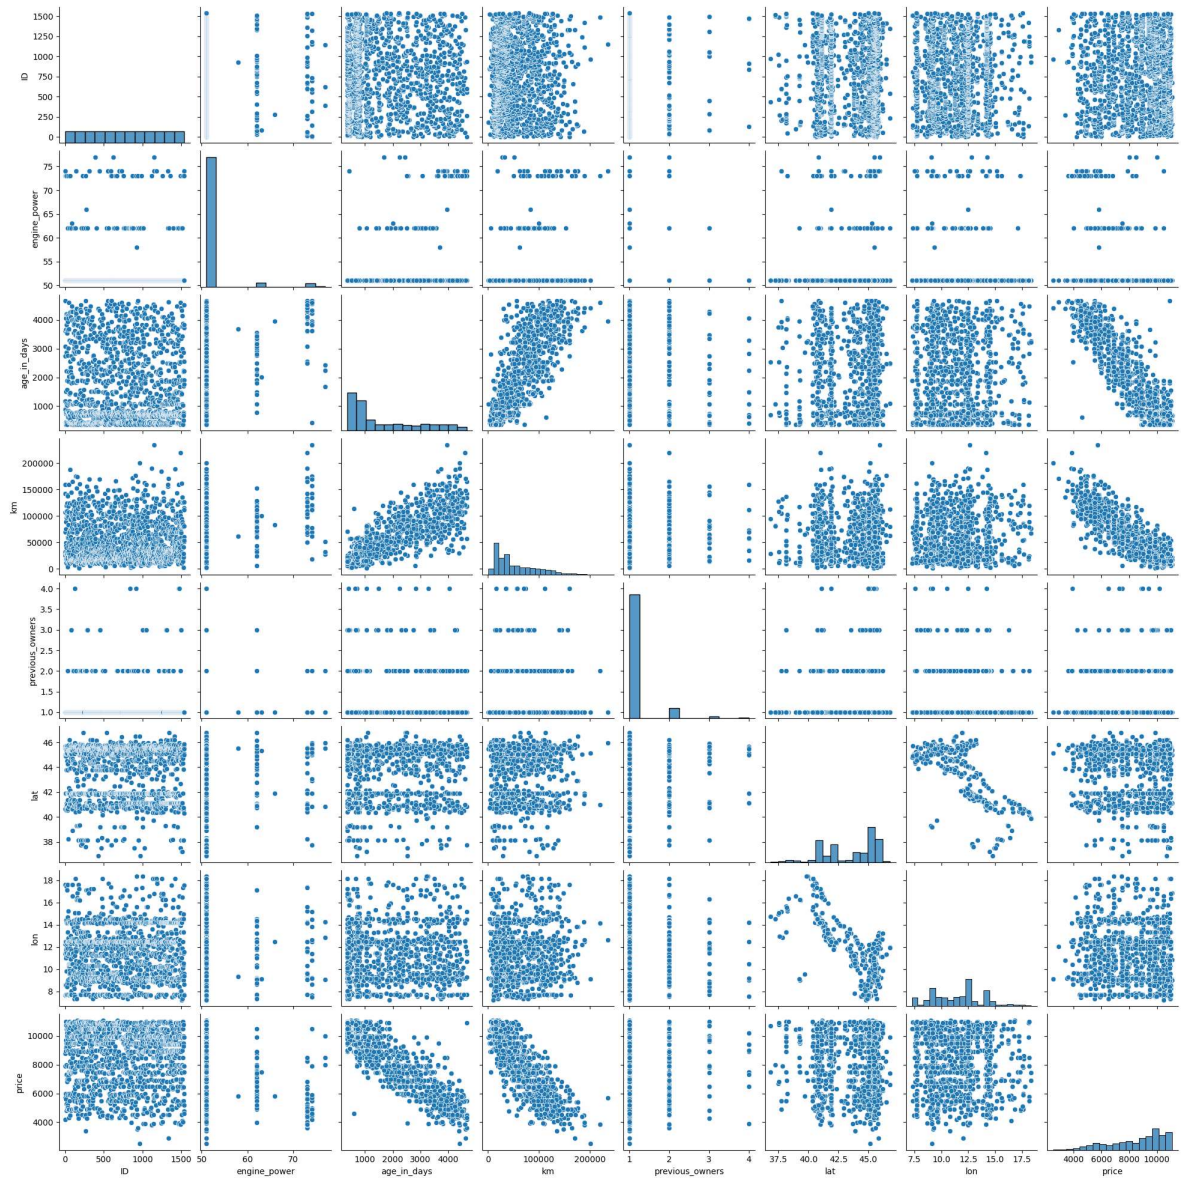
```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1538 entries, 0 to 1537
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  -
0   ID              1538 non-null   int64
1   model           1538 non-null   object
2   engine_power    1538 non-null   int64
3   age_in_days     1538 non-null   int64
4   km              1538 non-null   int64
5   previous_owners 1538 non-null   int64
6   lat             1538 non-null   float64
7   lon             1538 non-null   float64
8   price           1538 non-null   int64
dtypes: float64(2), int64(6), object(1)
memory usage: 108.3+ KB
```

In [65]: `df.columns`

```
Out[65]: Index(['ID', 'model', 'engine_power', 'age_in_days', 'km', 'previous_owner
s',
               'lat', 'lon', 'price'],
              dtype='object')
```


```
In [66]: sns.pairplot(df)
```

```
Out[66]: <seaborn.axisgrid.PairGrid at 0x16973151f50>
```



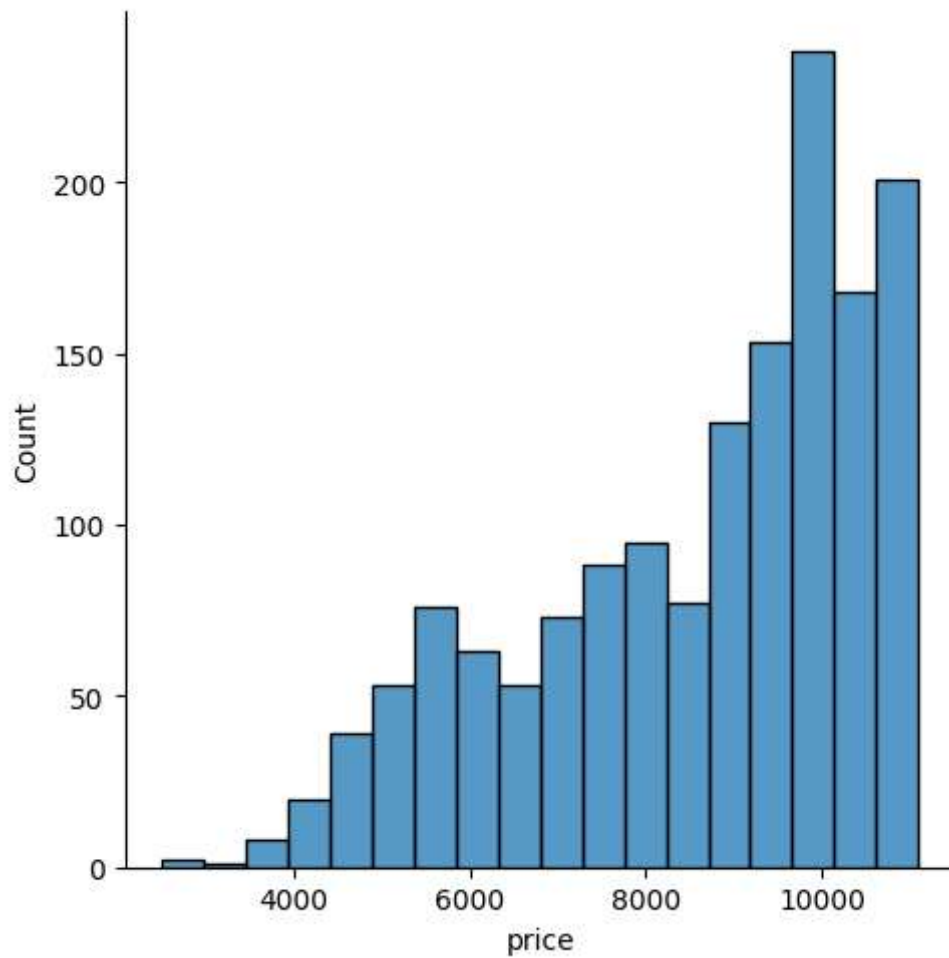
```
In [67]: sns.displot(['Price'])
```

```
Out[67]: <seaborn.axisgrid.FacetGrid at 0x1697334d190>
```




```
In [68]: sns.displot(df['price'])
```

```
Out[68]: <seaborn.axisgrid.FacetGrid at 0x1696f199c50>
```




```
In [69]: sns.displot(df['engine_power'])
```

```
Out[69]: <seaborn.axisgrid.FacetGrid at 0x169788e62d0>
```




```
In [70]: fiatdf=df[['ID', 'engine_power', 'age_in_days', 'km', 'previous_owners',  
                  'lat', 'lon']]  
sns.heatmap(fiatdf.corr())
```

Out[70]: <Axes: >



```
In [45]: X=fiatdf[['ID', 'engine_power', 'age_in_days', 'km', 'previous_owners','lat',  
                  y=df['price']
```

```
In [71]: from sklearn.model_selection import train_test_split  
X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.3,random_state=  
from sklearn.linear_model import LinearRegression  
regr=LinearRegression()  
regr.fit(X_train,y_train)  
print(regr.intercept_)
```

8971.19568349988


```
In [72]: coeff_df=pd.DataFrame(regr.coef_,X.columns,columns=['coefficient'])  
coeff_df
```

Out[72]:

	coefficient
ID	-0.046704
engine_power	11.646408
age_in_days	-0.898018
km	-0.017232
previous_owners	26.400886
lat	32.189709
lon	0.161073


```
In [73]: predictions=regr.predict(X_test)  
plt.scatter(y_test,predictions)
```

Out[73]: <matplotlib.collections.PathCollection at 0x1697abcf590>



```
In [74]: sns.displot((y_test,predictions),bins=50)
```

```
Out[74]: <seaborn.axisgrid.FacetGrid at 0x1697ab66a50>
```



```
In [75]: from sklearn import metrics
print('MAE:',metrics.mean_absolute_error(y_test,predictions))
print('MSE:',metrics.mean_squared_error(y_test,predictions))
print('MAE:',np.sqrt(metrics.mean_squared_error(y_test,predictions)))
```

MAE: 593.0876179519931

MSE: 551442.6799691801

MAE: 742.5918663500026

```
In [76]: #accuracy
regr=LinearRegression()
regr.fit(X_train,y_train)
regr.fit(X_train,y_train)
print(regr.score(X_test,y_test))
```

0.8597136704308868


```
In [77]: df.fillna(method='ffill',inplace=True)
```

```
In [78]: x=np.array(df['age_in_days']).reshape(-1,1)
y=np.array(df['km']).reshape(-1,1)
df.dropna(inplace=True)
```

```
In [79]: X_train,X_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
regr.fit(X_train,y_train)
regr.fit(X_train,y_train)
```

```
Out[79]: ▾ LinearRegression
LinearRegression()
```

```
In [80]: y_pred=regr.predict(X_test)
plt.scatter(X_test,y_test,color='y')
plt.plot(X_test,y_pred,color='b')
plt.show()
```



```
In [82]: #elasticnet
from sklearn.linear_model import ElasticNet
regr=ElasticNet()
regr.fit(x,y)
print(regr.coef_)
print(regr.intercept_)
y_pred_elastic=regr.predict(X_train)
mean_squared_error=np.mean((y_pred_elastic-y_train)**2)
print("Mean Squared Error on test set",mean_squared_error)
```

[25.89689696]

[10640.73996329]

Mean Squared Error on test set 2769977842.8600845

In []: